

## 1. Visualizing Loss Landscapes: PINNs vs. Data-Driven (40 points + 20 points)

Physics-Informed Neural Networks (PINNs) have emerged as a powerful tool for solving partial differential equations (PDEs). However, unlike traditional supervised learning (Data-Driven), PINNs often suffer from “pathological” optimization behaviors. The loss landscape of a PINN is determined by high-order derivatives of the network, which can result in highly non-convex, rough, and ill-conditioned landscapes, especially as the frequency of the solution increases.

In this exercise, you will investigate the spectral bias and optimization complexity of neural PDE solvers. You will solve a multiscale Poisson equation using both a **PINN** (residual-based) and a **Data-Driven** (supervised) approach. Your goal is to visualize and compare how the loss landscape degrades as the frequency of the target solution increases.

### Problem Statement

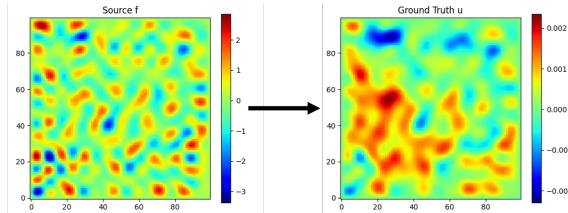
We focus on a prototypical linear elliptic PDE, the **Poisson Equation**, defined on a 2D square domain  $D = [0, 1]^2$ :

$$\begin{aligned} -\Delta u &= f, \quad \text{in } D, \\ u &= 0, \quad \text{on } \partial D. \end{aligned}$$

To control the complexity of the problem, we define a source term  $f$  comprised of  $K$  spatial scales. The source term and its corresponding analytical solution  $u(x, y)$  are given by:

$$\begin{aligned} f(x, y) &= \frac{\pi}{K^2} \sum_{i,j=1}^K a_{ij} \cdot (i^2 + j^2)^r \sin(\pi i x) \sin(\pi j y) \\ u(x, y) &= \frac{1}{\pi K^2} \sum_{i,j=1}^K a_{ij} \cdot (i^2 + j^2)^{r-1} \sin(\pi i x) \sin(\pi j y) \end{aligned}$$

where  $r = 0.5$  and coefficients  $a_{ij} \sim \mathcal{N}(0, 1)$ . The parameter  $K$  acts as an indicator for the problem’s complexity: as  $K$  increases, the solution contains higher frequency components.



**Figure 1:** Visualization of one data sample with  $K = 16$ .

### Tasks

#### Task 1: Data Generation (10 points)

Write a script to generate a dataset of input-output pairs  $(f^{(i)}, u^{(i)})$ . Discretize the domain  $D$  into an  $N \times N$  structural grid (e.g.  $N = 64$ ). Generate different samples by randomizing the coefficients

$a_{ij}$  for each sample. Plot 3 examples of the source field  $f$  and the corresponding solution field  $u$  for  $K = 1, 4, 8, 16$ .

### Task 2: Implementation and Training (30 points)

Implement a standard Multi-Layer Perceptron (MLP) with 3-4 hidden layers. The MLP here acts as a direct function approximator  $u_\theta : \mathbb{R}^2 \rightarrow \mathbb{R}$ . The input to the network is the spatial coordinates  $(x, y)$  of a point, and the output is the predicted scalar value  $\hat{u}$  at that location. For this task, you will optimize the network to solve the PDE for one specific source term  $f$  only. The model is expected to fit the solution for this single case; no generalization to other source terms or distributions is required in this exercise. You will perform this optimization twice: once using the Data-Driven loss and once using the PINN loss.

- **(20 points)** For PINN loss, you need to define the loss calculator similar to the following form:  $\mathcal{L}_{\text{PINN}}(\theta) = \frac{1}{N_f} \sum |-\Delta \hat{u}_\theta - f|^2 + \lambda \mathcal{L}_{\text{BC}}$ . Then use the automatic differentiation to calculate the gradient.
- **(10 points)** For Data-Driven loss, the loss calculator:  $\mathcal{L}_{\text{Data}}(\theta) = \frac{1}{N_d} \sum |\hat{u}_\theta - u_{\text{exact}}|^2$ .

You need to train both models for different complexity levels: **Low** ( $K = 1$ ), **Medium** ( $K = 4$ ), and **High** ( $K = 16$ ). The specific source term  $f$  to be solved can be defined by yourself. We recommend selecting one representative instance from the samples generated in the Task 1 to ensure consistency in your comparisons. Please record your final  $L_2$  relative error, training loss curve, and model prediction in your report.

**Note:** To accelerate convergence, adopt a combined optimizer strategy: start with **Adam** and switch to **L-BFGS** for fine-tuning.

### Task 3: Loss Landscape Visualization (BONUS - 20 points)

Using the techniques from *Li et al. (2018)*, visualize the geometry of the loss function around the converged (or stuck) solution parameters  $\theta^*$ . Compute the loss on a 2D plane:

$$g(\alpha, \beta) = \mathcal{L}(\theta^* + \alpha\delta + \beta\eta) \quad (1)$$

where  $\delta$  and  $\eta$  are two principal direction vectors.

- **(15 points):** Provide 2D contour plots or 3D surface plots of the landscapes for both PINN and Data-Driven models across the three  $K$  levels.
- **(5 points):** Qualitatively describe the differences. Is the PINN landscape sharper? Does it exhibit more local minima as  $K$  increases compared to the Data-Driven approach? Reference the “Failure modes” described by *Krishnapriyan et al. (2021)* in your discussion.

## References

- Krishnapriyan et al. (2021). Characterizing possible failure modes in physics-informed neural networks.
- Li et al. (2018). Visualizing the loss landscape of neural nets.